

Beyond Searching a Village: Learning to Recommend Diverse and Successful Collaborative Teams

Mahdis Saeedi

School of Computer Science
University of Windsor, Windsor, ON, Canada
msaeedi@uwindsor.ca

Hossein Fani

Faculty of Information
University of Toronto, Toronto, ON, Canada
hossein.fani@utoronto.ca

Abstract

Team recommendation involves selecting skilled experts to form an *almost surely* successful collaborative team, or refining the team composition to maintain and/or excel at performance. To address the tedious and error-prone manual process, various computational approaches have been proposed, especially for the web-scale social networks and widespread online collaboration and diversity of interactions. In this tutorial, with a brief overview of pioneering subgraph optimization approaches and their shortfalls, we deliberately focus on the recent learning-based approaches, with a particular in-depth exploration of graph neural network-based methods. More importantly, and for the first of its kind, we then discuss team *refinement*, which involves structural adjustments or expert replacements to enhance team performance in dynamic environments. Finally, we discuss training strategies, benchmarking datasets, and open-source libraries, along with future research directions and real-world applications. Further resources are available at <https://fani-lab.github.io/OpeNTF/tutorial/recsys26>.

CCS Concepts

• **Information systems** → **Recommender systems**; • **Human-centered computing** → **Social recommendation**; • **Computing methodologies** → **Neural networks**; • **Theory of computation** → **Graph algorithms analysis**.

Keywords

Team Recommendation; Social Information Retrieval; OpeNTF; Adila

ACM Reference Format:

Mahdis Saeedi and Hossein Fani. 2026. Beyond Searching a Village: Learning to Recommend Diverse and Successful Collaborative Teams. In *20th ACM Conference on Recommender Systems (RecSys '26)*, September 27–October 02, 2026, Minneapolis, MN, USA. ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/3773078.3831859>

1 Motivation

Team recommendation aims to automate the formation of collaborative teams of experts capable of successfully accomplishing a given task. Applications span a wide range of domains. In emergency response, it involves assembling ad hoc teams of diverse professionals while accounting for expertise, communication efficiency, availability, and logistical constraints. In scientific publishing, it supports the assignment of peer reviewers, where identifying an

```
[10 min] Pioneering Techniques
        [10 min] Subgraph Optimization Objectives and Techniques
[60 min] Learning Team Recommendation
[35 min] Team Refinement
        [15 min] Structural Adjustment
        [20 min] Membership Substitution
[20 min] Training Strategies
[10 min] Evaluation Methodologies
        [05 min] Datasets
        [05 min] Effectiveness/Efficiency
[15 min] Future Directions
        [05 min] Fair and Diverse Team Recommendation
        [05 min] Spatial Team Recommendation
        [05 min] Multi-Objective Neural Team Recommendation
[20 min] Real-World Applications
[10 min] OpeNTF and Adila
```

Figure 1: Our tutorial's timeline.

appropriate review panel remains a significant bottleneck due to the largely manual selection process. In education, educators seek to partition students into collaborative teams that promote peer learning, engagement, and social skill development, a need that has become increasingly important with the growth of large-scale online classes. On online freelance platforms, where abundant information is available about experts' skills, past performance, and collaboration history, team recommendation enables forming effective teams for numerous projects with evolving requirements.

Historically, collaborative teams were formed manually based on human experience and instinct; a tedious, error-prone, and sub-optimal process due to unfair biases, multiple criteria to optimize like budget, time and team size limitations, and a large number of candidates, among other reasons, which make it all the more difficult and impossible for a web-scale pool of experts. Hence, artificial intelligence and machine learning researchers have long been developing computational models learning through massive collections of experts and efficiently learning relationships between experts and their skills in the context of successful and *unsuccessful* teams and excel at recommending *almost surely* successful teams, and at helping existing teams maintain or even improve their success.

2 Prior Tutorials

Despite the abundance, there had been *no* comprehensive tutorial on team recommendation methods until recently when we began to fill the gap by tutorials at UMAP24¹ [16], centered on *online* skilled users, SIGIR-AP24² [17], based on a novel taxonomy from a computational perspective with a focus on real-world applications, and at WSDM25³ [18] and CIKM25⁴ [14], to review seminal solutions



This work is licensed under a Creative Commons Attribution 4.0 International License. RecSys '26, Minneapolis, MN, USA

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2284-4/2026/09
<https://doi.org/10.1145/3773078.3831859>

¹www.um.org/umap2024/tutorials/

²www.sigir-ap.org/sigir-ap-2024/tutorial/

³www.wsdm-conference.org/2025/tutorials/

⁴www.cikm-conference.org/2025/tutorials/




Figure 2: Taxonomy of team recommendation methods.

with a special focus on the emerging graph neural network-based methods. In this tutorial, from Figure 1 and 2, we expand on them based on our recently published [15] and ongoing surveys of neural approaches that leverage graph neural networks [13, 21], sequence-to-sequence models and transformers [20]. Notably, we further investigate team refinement methods, such as expert replacement through reinforcement learning in dynamic sport teams or online video games [11, 22]. We also present proposed methods to mitigate inherent popularity and gender biases via curriculum learning and loss regularization [3, 9, 12].

3 Targeted Audience

In this tutorial, (1) we target *beginner* or *intermediate* researchers, industrial practitioners with a broad interest in developing web-scale recommender systems, willing to have a whole picture of team recommendation techniques. (2) Furthermore, we target the graph neural network community for a comprehensive review of subgraph optimization objectives and call them for further development of effective yet efficient graph machine learning frameworks with a special focus on team recommendation. Last, this tutorial enables (3) organizations and practitioners to compare different approaches and readily pick the suitable one for their application to form or maintain successful teams. The target audience will be those *familiar* with graph theory and neural architectures. We provide sufficient details about advanced techniques so that the content becomes understandable to those with a fair knowledge of fundamentals.

4 Tutorial Summary (180 minutes)

We start our tutorial with a brief introduction to the pioneering graph-based team recommendation algorithms [7] based on a taxonomy of computational methods, as shown in Figure 2, then continue to explore the learning-based team recommendation [1, 3–6, 8, 20] and team refinement methods [2, 11, 22, 23], focusing on modern methods based on graph neural networks and reinforcement learning. We explain strategies to train team recommenders, including negative sampling [4], temporal training [6], and curriculum learning [3]. We also discuss the benchmark datasets, what has been

considered successful for a team as the ground truth, and metrics utilized to measure the quality of the recommended teams.

Toward the end, we discuss open issues and future directions, including:

- **Fair and Diverse Teams:** Existing methods heavily influenced by hidden *unfair* societal biases [9, 10, 12]. Meanwhile, social science research provides compelling evidence about the synergistic effects of diversity on team performance. In our tutorial, we introduce notions of fairness and protected attributes and study debiasing algorithms to mitigate the potential unfairness in the models' recommendations.
- **Spatial Team Recommendation:** In search of an optimal team, companies further look for experts in a region where the organization is *geographically* based. We bring forth the *spatial* team recommendation problem; that is, given a set of experts, skills and geolocations, the goal is to determine if the combination of skills and geolocations in forming teams has synergistic effects.
- **Multi-Objective Optimization:** In real-world scenarios, balancing multiple, often conflicting objectives (e.g., team effectiveness and experts' workload distribution) requires a training process guided by a loss function that explicitly accounts for multiple objectives. However, existing approaches commonly frame the problem as a classification or link prediction task, aim to maximize the coverage of the required skills, and overlook designing a task-aware loss function.

Moreover, we explain novel applications of team recommendation, including learning group recommendation, peer-review assignments, and palliative care [19].

Finally, we present OpeNTF⁵ [5, 8], a benchmark library for neural models that can efficiently preprocess large-scale datasets, can be easily extended or customized to new models, and is extensible to new datasets from other domains. We also introduce Adila⁶ [9], which enables further post-processing reranking to the list of recommended experts to ensure a fair outcome, and vivaFemme [12], an in-processing method for female-advocate loss regularization during model training.

⁵github.com/fani-lab/OpeNTF

⁶github.com/fani-lab/Adila

